

Spark Assignment

MSc in Business Analytics

Big Data Systems

Deadline: Fri April 1st 2022, 23:59

Background

You have been hired by a small bookstore company that wants to use data science techniques to optimize their sales. It has been assigned to you to analyse a dataset of books metadata using Apache Spark (and PySpark, in particular) to reveal useful insights. You can find the dataset in the e-class page.

Task 1 [50 points]

Your first task is to explore the dataset. You need to use SparkSQL with Dataframes in a Jupyter notebook that delivers the following:

- It uses the `json()` function to load the dataset.
- It counts and displays the number of books in the database.
- It counts and displays the number of e-books in the database (based on the “is_ebook” field).
- It uses the `summary()` command to display basic statistics about the “average_rating” field.
- It uses the `groupby()` and `count()` commands to display all distinct values in the “format” field and their number of appearances

Your deliverable should be a ready-to-run Jupyter notebook (named “id-t1.ipynb”), containing your details (name, id) and explanations for each step of the code.

Task 2 [30 points]

For this task you continue to work with SparkSQL. This time, you need to provide a Jupyter notebook (again using PySpark and Dataframes) that delivers the following:

- It returns the “book_id” and “title” of the book with the largest “average_rating” that its title starts with the first letter of your last name.
- It returns the average “average_rating” of the books that their title starts with the *second* letter of your last name.
- It returns the “book_id” and “title” of the Paperback book with the most pages, when only books with title starting with the *third* letter of your last name are considered.

Your deliverable should be a ready-to-run Jupyter notebook (named “id-t2.ipynb”), containing your details (name, id) and explanations for each step of the code.

Task 3 [20 points]

As a final task, your supervisor assigned to you to investigate if it is possible to train a linear regression model (using `LinearRegression()` function) that could predict the “average_rating” of a book, using as input, its “language_code”, its “num_pages”, its “ratings_count”, and its “publication_year”. Again you should use Python and Dataframes, this time with MLlib. You should pay attention to transform the string-based input features (“language_code” and “publication_year”) using the proper representation format, and you should explain your choices. Your code should (a) prepare the feature vectors, (b) prepare the training and testing datasets (70%-30%), (c) train the model, and (d) evaluate the accuracy of the model (based on the Rsquared metric) and display the corresponding metric on the screen.

Your deliverable should be a ready-to-run Jupyter notebook (named "id-t3.ipynb"), containing your details (name, id) and explanations for each step of the code.

Submission instructions & honor code

Your code files should be fully replicable and readable (documentation comments are required and appreciated). Your code should work with Spark v.3 (using PySpark) and should be ready to be executed (e.g., containing all the required import statements). Code failing to execute or producing wrong results will be penalised. You understand that this is an individual assignment, and as such you must carry it out alone. You may discuss with your fellow students to better understand the tasks/questions but you should not ask them to share their answers with you or to help you by giving you specific advice.

GOOD LUCK!